

www.locuz.com

Distributed Computing with MPI

Presented By: Mandeep Kumar




Converge to the Cloud

Agenda

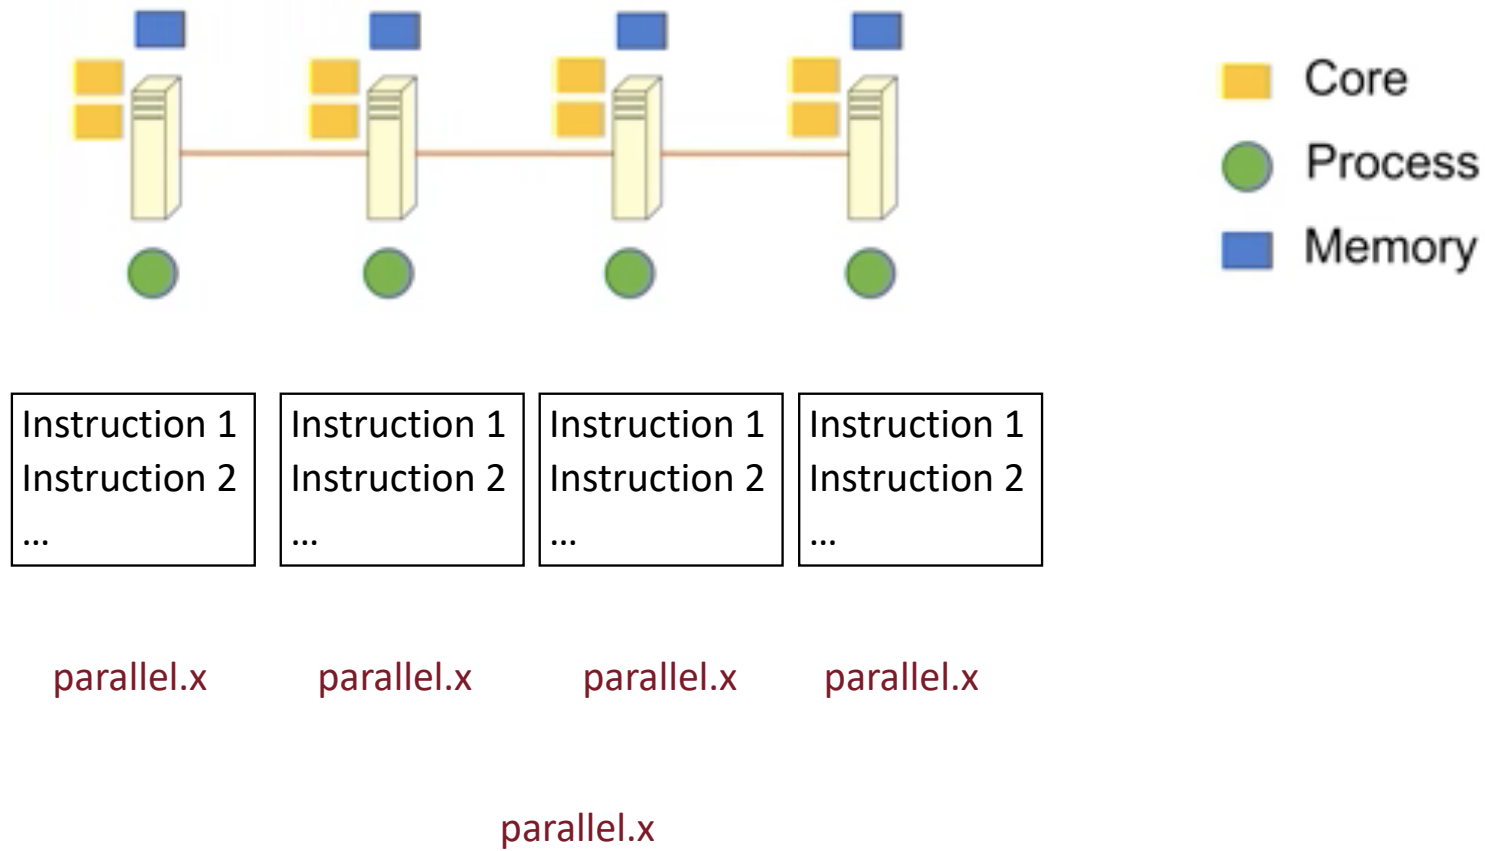
- **MPI Introduction**
- **Communicator and communication**
- **Point to point communication**
- **Collective communication**

System Details



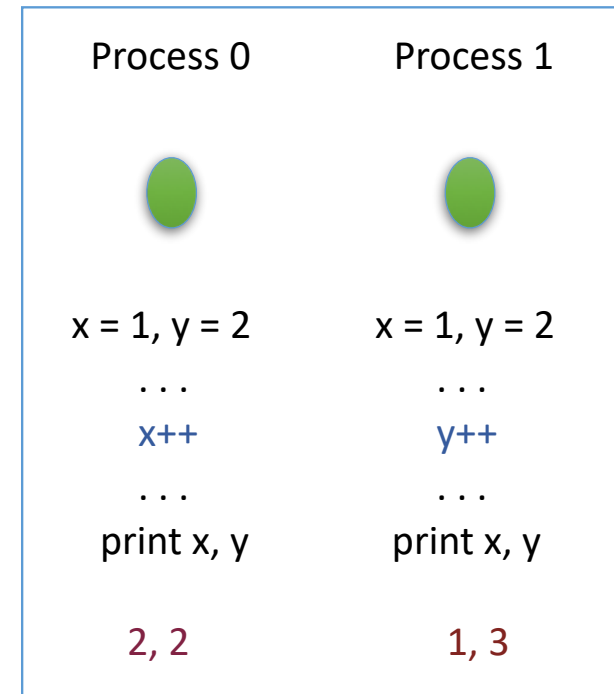
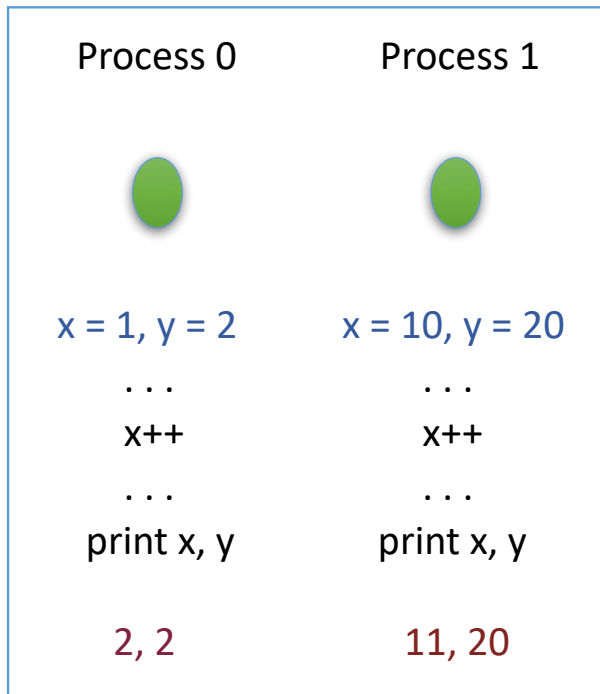
Distributed Memory (each process has its own address space)

Parallel Execution



Distinct Process Address Space

Program
Order



What is MPI?

- MPI stands for Message Passing Interface and is a library specification for message-passing, proposed as a standard by a broadly based committee of vendor, implementors, and users.
- API for distributed- memory programming
 - parallel code that runs across multiple computers (nodes)
- MPI consists of
 - a header file `mpi.h`
 - a library of routines and functions, and
 - a runtime system
- MPI is for parallel computers, clusters, and heterogeneous networks.
- Use from C/C++, Fortran, Python, R, ...
- More than 200 routines
- Using only 10 routines are enough in many cases
 - Problem dependent

Example MPI routines

The following routines are found in nearly every program that uses MPI

- `MPI_Init()` starts the MPI runtime environment.
- `MPI_Finalize()` shuts down the MPI runtime environment.
- `MPI_Comm_size()` gets the number of processes, N_p .
- `MPI_Comm_rank()` gets the process ID of the current process which is between 0 and $N_p - 1$, inclusive.

(These last two routines are typically called right after `MPI_Init()`.)

Parallel hello to you!

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char *argv[])
{
    // initialize MPI
    MPI_Init (&argc, &argv);

    printf ("Parallel hello to you!\n");

    // done with MPI
    MPI_Finalize();
}
```

Initialization

Finalization

MPI_Init

- gather information about the parallel job
- set up internal library state
- prepare for communication

MPI Code Execution Steps

- **Compile**

```
mpicc -o program.x program.c
```

- **Execute**

- `mpirun -np 1 ./program.x` (or `mpiexec` in place of `mpirun`)
 - Run 1 process on launch node
- `mpirun -np 6 ./program.x`
 - Run 6 process on launch node
- `mpirun -np 6 -f hostfile ./program.x`
 - Run 6 process on the nodes specified in the hostfile

<hostfile>


Node1:2

Node2:2

Node3:2

...

Multiple Processes on Multiple Cores and Nodes




```
mpirun -np 8 -f hostfile ./program.x
```

Communicator

- Communication handle among a group/collection of processes
 - processes are numbered 0,1,... to N-1
- Default Communicator:
 - `MPI_COMM_WORLD`
 - contains all processes
- Contains a mapping from MPI process ranks to processor ids
- Memory proportional to #processes in the group
- Query functions:
 - How many processes in total?
`MPI_Comm_size(MPI_COMM_WORLD, &nproc)`
 - What is my process ID?
`MPI_Comm_rank(MPI_COMM_WORLD, &rank)`

MPI_COMM_WORLD



Required in every MPI communication

Process Identification

- **MPI_Comm_size**: get the total number of process
- **MPI_Comm_rank**: get my rank

```
#include <stdio.h>
#include "mpi.h"
int main(int argc, char *argv[])
{
    int numtasks, rank, len;
    char hostname[MPI_MAX_PROCESSOR_NAME];

    // initialize MPI
    MPI_Init (&argc, &argv);

    // get number of tasks
    MPI_Comm_size (MPI_COMM_WORLD, &numtasks);

    // get my rank
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);


    // this one is obvious
    MPI_Get_processor_name (hostname, &len);

    printf ("Number of tasks=%d My rank=%d Running on %s\n", numtasks, rank, hostname);
    // done with MPI
    MPI_Finalize();
}
```

Rank of process

Total number of
processes

Multiple Processes on Multiple Cores and Nodes



`mpirun -np 8 -f hostfile ./program.x`

Node1:2
Node2:2
Node3:2
Node4:2



Note:

- All MPI processes (normally) run the same executable
- Each MPI process knows which rank it is
- Each MPI process knows how many processes are part of same job
- The processes run in a non-deterministic order

Sum of Squares of N Numbers

Serial

```
for i = 1 to N  
  sum += a[i] * a[i]
```



Parallel

```
for i = 1 to N/P  
  sum += a[i] * a[i]
```



```
for i = 1 to N/P  
  sum += a[i] * a[i]
```

```
for i = 1 to N/P  
  sum += a[i] * a[i]
```

```
for i = 1 to N/P  
  sum += a[i] * a[i]
```

```
for i = 1 to N/P  
  sum += a[i] * a[i]
```

```
for i = N/P * rank; i < N/P * (rank+1); i++  
  localsum += a[i] * a[i]  
Collect localsum, add up at one of the ranks
```


Parallel Sum of Array

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include "mpi.h"
#define N 100000000

int main(int argc, char *argv[])
{
    int numtasks, rank, len, rc, i, sidx;
    double etime, stime, a[N], value, localsum;
    char hostname[MPI_MAX_PROCESSOR_NAME];

    // initialize MPI
    MPI_Init (&argc, &argv);

    // get number of tasks
    MPI_Comm_size (MPI_COMM_WORLD, &numtasks);

    // get my rank
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
```

```
    MPI_Get_processor_name (hostname, &len);

    // random initialization
    srand(time(NULL));
    value = abs((numtasks-rank)*(20210401-rand()))%N;
    sidx = rank*N/numtasks;
    for (i=sidx; i<sidx+N/numtasks ; i++)
        a[i] = value;

    // compute local sum
    localsum=0.0;
    stime = MPI_Wtime();
    for (i=sidx; i<sidx+N/numtasks ; i++)
        localsum += a[i];
    etime = MPI_Wtime();

    printf ("%d: Time to sum: %lf\n", rank, etime - stime);
    // done with MPI
    MPI_Finalize();
}
```

Scalability Bottleneck:
Memory Bound


Performance

- Speedup

$$S_p = \frac{\text{Time (1 processor)}}{\text{Time (P processors)}}$$

- Efficiency


$$E_p = \frac{S_p}{P}$$



MPI Communication Types

- Point-to-point
- Collective

Point-to-Point Communication



MPI_Send

- Send data to another process

MPI_Send (buffer, count, datatype, dest, tag, comm)

| Arguments | Meanings |
|-----------|---------------------------------------|
| buf | starting address of send buffer |
| count | # of elements |
| datatype | data type of each send buffer element |
| dest | process ID (rank) destination |
| tag | message tag |
| comm | communicator |

- Examples:

C/C++: **MPI_Send**(&x, 1, MPI_INT, 5 , 0, MPI_COMM_WORLD);
Fortran: **MPI_Send**(x, 1, MPI_INTEGER, 5, 0, MPI_COMM_WORLD, ierr)

MPI_Recv

- Receive data from another process

MPI_Recv (buffer, count, datatype, src, tag, comm, status)

| Arguments | Meanings |
|-----------|---|
| buf | starting address of send buffer |
| count | # of elements |
| datatype | data type of each send buffer element |
| dest | process ID (rank) source |
| tag | message tag |
| comm | communicator |
| status | status object (an integer array in Fortran) |

- Examples:

```
C/C++: MPI_Recv(&x, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);  
Fortran: MPI_Recv(x, 1, MPI_INTEGER, 1, 0, MPI_COMM_WORLD, status, ierr)
```

Notes on MPI_Recv

- A message is received when the followings are matched:
 - Source (sending process ID/rank)
 - Tag
 - Communicator (e.g. MPI_COMM_WORLD)
- Wildcard values may be used:
 - MPI_ANY_TAG
(Don't care what the tag value is)
 - MPI_ANY_SOURCE
(Don't care where it comes from; always receive)

Send/Recv Example (C)

- Send an integer array `f[N]` from process 0 to process 1

```
int f[N], src=0, dest=1;
MPI_Status status;
// ...
MPI_Comm_rank( MPI_COMM_WORLD, &rank);

if (rank == src)           // process "dest" ignores this
MPI_Send(f, N, MPI_INT, dest, 0, MPI_COMM_WORLD);

if (rank == dest)         // process "src" ignores this
MPI_Recv(f, N, MPI_INT, src, 0, MPI_COMM_WORLD, &status);

//...
```


Send/Recv Example (F90)

- Send an integer array f(1:N) from process 0 to process 1

```
integer f(N), status(MPI_STATUS_SIZE), rank, src=0, dest=1, ierr
// ...
call MPI_Comm_rank( MPI_COMM_WORLD, rank, ierr);

if (rank == src) then                                !process "dest" ignores this
call MPI_Send(f, N, MPI_INT, dest, 0, MPI_COMM_WORLD, ierr)
end if

if (rank == dest) then                                !process "src" ignores this
call MPI_Recv(f, N, MPI_INT, src, 0, MPI_COMM_WORLD, status, ierr)
end if

//...
```

Send/Recv Example (cont'd)

- Before

| process 0 (send) | process 1 (recv) |
|------------------|------------------|
| f[0]=0 | f[0]=0 |
| f[1]=1 | f[1]=0 |
| f[2]=2 | f[2]=0 |

- After

| process 0 (send) | process 1 (recv) |
|------------------|------------------|
| f[0]=0 | f[0]=0 |
| f[1]=1 | f[1]=1 |
| f[2]=2 | f[2]=2 |

Blocking

- Function call does not return until the communication is complete MPI_Send and MPI_Recv are blocking calls
- Calling order matters
 - it is possible to wait indefinitely, called “deadlock”
 - improper ordering results in serialization (loss of performance)

Deadlock

- This code always works:

```
MPI_Comm_rank(comm, &rank);

if (rank == 0) {
    MPI_Send(sendbuf, count, MPI_INT, 1, tag, comm);
    MPI_Recv(recvbuf, count, MPI_INT, 1, tag, comm, &stat);
} else { // rank==1
    MPI_Recv(recvbuf, count, MPI_INT, 0, tag, comm, &stat);
    MPI_Send(sendbuf, count, MPI_INT, 0, tag, comm);
}
```

Deadlock

- This code deadlocks:

```
MPI_Comm_rank(comm, &rank);

if (rank == 0) {
    MPI_Recv(recvbuf, count, MPI_INT, 1, tag, comm, &stat);
    MPI_Send(sendbuf, count, MPI_INT, 1, tag, comm);
} else { // rank==1
    MPI_Recv(recvbuf, count, MPI_INT, 0, tag, comm, &stat);
    MPI_Send(sendbuf, count, MPI_INT, 0, tag, comm);
}
```

reason: MPI_Recv on process 0 waits indefinitely and never returns.

Non-blocking

- Function call returns immediately, without completing data transfer
 - Only “starts” the communication (without finishing)
 - MPI_Isend and MPI_Irecv
 - Need an additional mechanism to ensure transfer completion (MPI_Wait)
- Avoid deadlock
- Possibly higher performance
- Example: MPI_Isend & MPI_Irecv

MPI_Isend

```
MPI_Isend (buffer, count, datatype, dest, tag, comm, request)
```

- Similar to MPI_Send, except the last argument “request”
- Typical usage:

```
MPI_Request request_X, request_Y;  
MPI_Isend(..., &request_X);  
MPI_Isend(..., &request_Y);  
  
// ... some ground-breaking computations ...  
  
MPI_Wait(&request_X, ...);  
MPI_Wait(&request_Y, ...);
```

MPI_Irecv

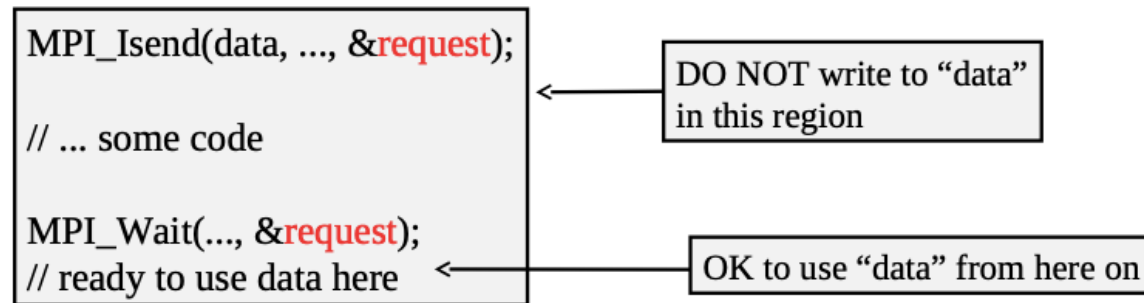
```
MPI_Irecv (buffer, count, datatype, dest, tag, comm, request)
```

- Similar to MPI_Recv, except the last argument “request”
- Typical usage:

```
MPI_Request request_X, request_Y;  
MPI_Irecv(..., &request_X);  
MPI_Irecv(..., &request_Y);  
  
// ... more ground-breaking computations ...  
  
MPI_Wait(&request_X, ...);  
MPI_Wait(&request_Y, ...);
```


Caution about MPI_Isend and MPI_Irecv

- The sending process should not access the send buffer until the send completes



MPI_Wait

```
MPI_Wait (MPI_Request, MPI_Status)
```

- Wait for an MPI_Isend/recv to complete
- Use the same “request” used in an earlier MPI_Isend or MPI_Irecv
- If they are multiple requests, one can use
MPI_Waitall(count, request[], status[]);
request[] and status[] are arrays.

Other variants of MPI Send/Recv

- MPI_Sendrecv
 - send and receive in one call
- Mixing blocking and non-blocking calls
 - e.g. MPI_Isend + MPI_Recv
- MPI_Bsend
 - buffered send
- MPI_Ibsend
- ... (see MPI standard for more)

Synchronization (MPI_Barrier)

- MPI_Barrier (comm)
- Every rank needs to call this function (for true synchronization)
- Caller returns only after all processes have entered the call

```
printf("Before barrier");  
MPI_Barrier (MPI_COMM_WORLD);  
printf("After barrier");
```

Collective Communication

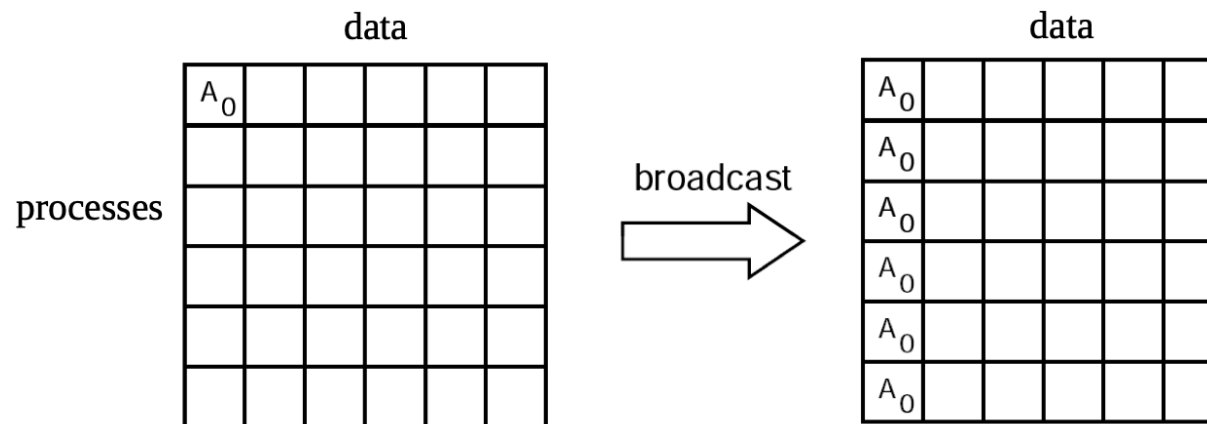
- One to all
 - MPI_Bcast, MPI_Scatter
- All to one
 - MPI_Reduce, MPI_Gather
- All to all
 - MPI_Alltoall

Implicit Synchronization

MPI_Bcast

`MPI_Bcast` (buffer, count, datatype, root, comm)

- Broadcasts a message from “root” process to all other processes in the same communicator



MPI_Bcast Example

- Broadcast 100 integers from process “3” to all other processes

C/C++

```
MPI_Comm comm;  
int array[100];  
// ...  
MPI_Bcast(array, 100, MPI_INT, 3, comm);
```

Fortran

```
INTEGER comm  
integer array(100)  
// ...  
call MPI_Bcast(array, 100, MPI_INTEGER, 3, comm, ierr)
```

MPI_Bcast vs MPI_Send+MPI_Recv

```
if (world_rank == root) {  
    // If we are the root process, send our data to everyone  
    for (int i = 0; i < world_size; i++) {  
        if (i != world_rank) { MPI_Send(data, count, datatype, i, 0, communicator);}  
    }  
} else { // if we are a receiver process, receive the data from the root  
    MPI_Recv(data, count, datatype, root, 0, communicator, MPI_STATUS_IGNORE);  
}
```

Vs

```
MPI_Bcast(data, num_elements, MPI_INT, 0, MPI_COMM_WORLD);
```

A simple way to speedup:

- Use tree based communication

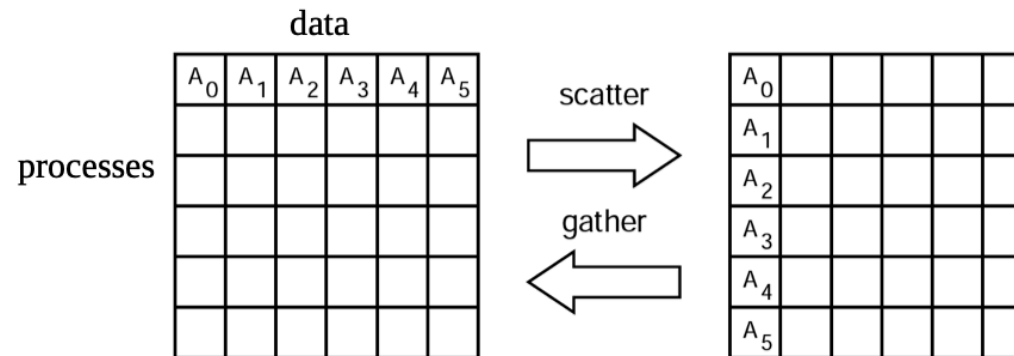
MPI_Gather & MPI_Scatter

- Gathers values from all processes to a root process

`MPI_Gather` (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvttype, root, comm)

- Scatters values to all processes from a root process


`MPI_Scatter` (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvttype, root, comm)



- When gathering, make sure the root process has big enough memory to hold the data (especially when you scale up the problem size).

Non blocking versions: `MPI_Igather` and `MPI_Iscatter`

MPI_Gather Example



```
MPI_Comm comm;
int np, myid, sendarray[N], root;
double *rbuf;
MPI_Comm_size( comm, &np);      // # of processes
MPI_Comm_rank( comm, &myid);    // process ID
if (myid == root)                // allocate space on process root
    rbuf = new double [np*N];

MPI_Gather(sendarray, N, MPI_INT, rbuf, N, MPI_INT, root, comm);
```

Scatter Gather Example

```
if (world_rank == 0) {rand_nums = create_rand_nums(elements_per_proc * world_size);}

// Create a buffer that will hold a subset of the random numbers
float *sub_rand_nums = malloc(sizeof(float) * elements_per_proc);

// Scatter the random numbers to all processes
MPI_Scatter(rand_nums, elements_per_proc, MPI_FLOAT, sub_rand_nums, elements_per_proc, MPI_FLOAT,
0, MPI_COMM_WORLD);

// Compute the average of your subset
float sub_avg = compute_avg(sub_rand_nums, elements_per_proc);

// Gather all partial averages down to the root process
float *sub_avgs = NULL;

if ( world_rank == 0) {sub_avgs = malloc(sizeof(float) * world_size);}

MPI_Gather(&sub_avg, 1, MPI_FLOAT, sub_avgs, 1, MPI_FLOAT, 0, MPI_COMM_WORLD);

// Compute the total average of all numbers
if (world_rank == 0) {float avg = compute_avg(sub_avgs, world_size);}
```


Variations of MPI_Gather/Scatter

- Variable data size
 - MPI_Gatherv
 - MPI_Scatterv
- Gather + broadcast (in one call)
 - MPI_Allgather
 - All processes send same amount of data
 - MPI_Allgatherv
 - Processes may send variable amount of data

MPI_Alltoall

`MPI_Alltoall` (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)

The j-th block sendbuf from process i is received by process j and is placed in the i-th block of recvbuf:



MPI_Reduce

```
MPI_Reduce (sendbuf, recvbuf, count, datatype, op, root, comm)
```

- Apply operation op to sendbuf from all processes and return result in the recvbuf on process “root”.
- op: MIN, MAX, SUM, PROD, ...
- Example:

MPI_Reduce (... , MPI_MAX, ...)

Output: 21

MPI_Reduce (... , MPI_MIN, ...).

Output: 1

21

5

1

8

3

2

13

MPI Resources on the Net

1. MPI Tutorial

Introduction, background, and basic information. Topics include MPI Environment Management, Point-to-Point Communications, and Collective Communications routines
<https://computing.llnl.gov/tutorials/mpi/>

2. Laplace Exercise

Serial and parallel exercise, domain decomposition, C and Fortran template codes.

<https://docplayer.net/61527559-Laplace-exercise-john-urbanic-parallel-computing-scientist-pittsburgh-supercomputing-center-copyright-2017.html>

3. Home page of Prof. William D. Gropp

<http://wgropp.cs.illinois.edu/>

4. Book: Using MPI- Portable Parallel Programming with the Message-Passing Interface by William Gropp, Ewing Lusk and Anthony Skjellum (MIT Press, Third edition)
<https://ieeexplore.ieee.org/book/6267273>

5. Book: Introduction to Parallel Computing by Gramma, Gupta, Karypis and Kumar (Pearson)

<https://www-users.cs.umn.edu/~karypis/parbook/>

6. Book: Parallel Programming in C with MPI and OpenMP by Michale J Quinn (McGrawhill)

<https://www.amazon.in/Parallel-Programming-C-MPI-OpenMP/dp/0072822562>

7. Book: Iterative Methods for Sparse Linear Systems by Yousef Saad

https://www-users.cs.umn.edu/~saad/IterMethBook_2ndEd.pdf

8. NPTEL Course on Matrix Solvers- Somnath Roy

Topic covered are Direct solvers, Iterative solvers, Krylov methods, Preconditioners, Domain Decomposition and Multigrid Methods
<https://nptel.ac.in/courses/111/105/111105111/>

Thanks!

